

TOWARDS GLOBAL GEOMETRY PRESERVATION IN TEXT EMBEDDING DISTILLATION

Anonymous authors

Paper under double-blind review

ABSTRACT

A text embedding teacher stores knowledge in the geometry it induces over an entire corpus, but supervising all pairwise relations is infeasible. Standard relational knowledge distillation (KD) instead forms its supervision support from examples that happen to co-occur in a mini-batch, so relation exposure is independent of teacher relevance: an important teacher relation and an unrelated corpus pair are equally likely to be supervised. We argue that a corpus-level relational objective should separate two decisions that mini-batch KD conflates—which relations are *eligible* for supervision, and how eligible relations are *realized* efficiently—and that preserving global geometry depends on both alignment, reproducing the supervised relations accurately, and coverage, supervising most of the relational mass the teacher considers important. We introduce **Global Geometry Preservation Knowledge Distillation (GGPKD)**, which lets teacher geometry choose the support and uses mini-batches only to realize it. GGPKD encodes the corpus once with the teacher, builds a k -nearest-neighbor graph whose per-row bandwidth is read off the retrieved similarities, and supervises each anchor on its own retrieved neighborhood, every column included. The deduplicated candidate pool of a batch is encoded once and reused twice more: to calibrate similarity levels across the batch, and to supervise additional teacher-selected rows without any extra student encoding. Supports, bandwidths, temperatures, and term weights are all derived from the graph, leaving two hyperparameters: the graph width and the auxiliary-row weight. Across nine standard sentence-embedding benchmarks and three teacher–student settings, GGPKD obtains the best mean score in every setting, with the largest and most consistent gains on semantic textual similarity. GGPKD requires no task labels and no online teacher inference, and only the student encoder is used at deployment.

1 INTRODUCTION

Text embedding models have become a fundamental component of modern NLP systems, supporting applications such as semantic search, retrieval, clustering, semantic textual similarity, and classification (Gao et al., 2023; Muennighoff et al., 2023). While larger embedding models can achieve strong performance, their computational and memory requirements may limit practical deployment, motivating the development of compact models that retain high-quality representations (Zhang et al., 2025). While larger text embedding models often achieve stronger performance on benchmark such as MTEB (Muennighoff et al., 2023), their increased parameter counts and representation dimensionality can incur substantial computational, memory, and deployment costs.

Knowledge distillation (KD) offers a practical way to reduce this deployment burden by transferring knowledge from a high-capacity teacher to a more compact student (Hinton et al., 2015). Recent embedding distillation frameworks have demonstrated strong empirical performance through direct teacher–student representation alignment (Vujanic & Rückstieß, 2026). For text embedding models, however, relational structure is also important, as many downstream applications operate directly on the similarities, distances, and relative ordering induced among embedded texts (Reimers & Gurevych, 2019; Muennighoff et al., 2023). Pointwise representation alignment does not explicitly constrain these inter-text relationships and may therefore leave important aspects of the teacher’s embedding geometry uncontrolled. This has motivated relational and structure-preserving distillation methods that explicitly transfer teacher-induced similarities, distances, angles, and other geometric

relationships (Tung & Mori, 2019; Park et al., 2019; Kim et al., 2023; Zhang et al., 2024; Dao et al., 2026).

At the corpus level, relational distillation would ideally account for the geometry induced among all training texts. Directly realizing such an objective, however, requires a quadratic number of pairwise relations. Relational KD is therefore commonly instantiated within mini-batches, where only relations among co-occurring texts are available for supervision (Tung & Mori, 2019; Park et al., 2019; Passalis et al., 2021). Prior work has identified this batch locality as a limitation: batch-wise objectives observe only a fraction of dataset-level relations, with relational coverage becoming increasingly sparse as the dataset grows (Qian et al., 2022; Jung et al., 2023). Full-kernel approximations and memory-augmented methods have consequently been explored to extend relational supervision beyond individual batches (Qian et al., 2022; Jung et al., 2023; Yang et al., 2022b). These approaches broaden relational supervision beyond individual batches, but do not address which corpus relations should be prioritized under a limited relational budget. We formulate this as a support-selection problem: mini-batch KD lets batch co-occurrence decide which relations are supervised at all, whereas GGPKD lets teacher geometry decide which relations are eligible and uses mini-batches only to realize them.

We argue that this approximation introduces an under-emphasized bottleneck. Relational KD involves two distinct decisions: *which relations should be supervised and how those relations should be computed efficiently*. Mini-batch KD entangles the two. Although the relations themselves are defined by the teacher, their supervision support is determined by random batch co-occurrence. Consequently, an important teacher relation and an irrelevant corpus pair have the same probability of being exposed during training. Optimizing longer over such batches does not change which parts of the teacher geometry are observed, only how often the same uninformed sample is revisited. This observation suggests a different principle: *teacher geometry should determine which relations are eligible for supervision, while mini-batches should serve to realize those relations efficiently*.

Based on this principle, we introduce GGPKD, a framework for teacher-selected relational distillation. GGPKD encodes the corpus once with the teacher and builds a sparse nearest-neighbor graph from the cached embeddings. Each row of this graph is turned into a transition distribution whose bandwidth is read off the retrieved similarities rather than tuned, which makes the target invariant to how a particular teacher scales its cosines. An anchor is then supervised on exactly the columns its own row puts mass on, so the supervision support is a deterministic function of teacher geometry and is identical in every epoch. Mini-batches enter only as a computational device: the candidate sets of a batch are deduplicated into a shared pool that is encoded once, and that single encoding pass supports three uses of the same vectors—matching each anchor’s transition row, calibrating similarity levels across the pool, and supervising the teacher-approved rows of non-anchor pool members. Batch composition therefore determines which eligible relations can be realized in the current update, not which relations are considered relevant. Because every bandwidth, temperature, and term weight is derived from the graph, the method exposes two hyperparameters: the graph width and the weight of the auxiliary-row term.

Across nine datasets and three teacher–student settings, GGPKD attains the best mean score in all three settings. The gains are largest and most consistent on semantic textual similarity, while pair-classification results are less uniform. Our contributions are:

- We formulate relational knowledge distillation as a *support-selection* problem, and separate relation eligibility, which should be decided by teacher geometry, from relation realization, which is a computational concern. Under this view, conventional mini-batch training selects relational supervision independently of teacher relevance, and corpus-level geometry preservation is controlled by two factors: alignment on the supervised relations and coverage of the teacher’s relational mass.
- We introduce GGPKD, which realizes this principle with a cached k NN teacher graph, per-row bandwidths derived from the retrieved similarities, deterministic row-shaped supports, and amortized auxiliary-row supervision that reuses already-encoded pool members. Every temperature and weight is derived from the graph, leaving two hyperparameters.
- We evaluate on nine datasets across three teacher–student settings, where GGPKD obtains the best mean score in every setting and the strongest results on semantic textual similarity,

and show that its two hyperparameters are flat enough that the defaults are not selected on test results.

2 RELATED WORK

2.1 TEXT EMBEDDING DISTILLATION

Knowledge distillation transfers behavior from a large teacher to a compact student (Hinton et al., 2015). Language-model compression commonly matches output distributions, hidden states, or attention maps (Sanh et al., 2019; Jiao et al., 2020; Wang et al., 2020; Xu et al., 2024), whereas embedding-specific methods target quantities used by similarity-based applications. Sparse Logit Sampling shows that deterministic top- K teacher logits induce truncation bias and uses randomized teacher-logit sampling to preserve the full teacher distribution in expectation (Anshumann et al., 2025). In that setting, however, the student still produces a full vocabulary vector; in corpus-level embedding KD, scoring an additional relation requires encoding another text, so support selection controls both approximation and encoder cost.

DistilCSE combines teacher supervision with contrastive learning (Gao et al., 2021; Xu et al., 2023); EmbedDistill transfers relative query–document geometry (Kim et al., 2023); and recent systems align embeddings, layers, tokens, or heterogeneous representation spaces (Lee et al., 2024; Zhang et al., 2024; Vujanic & Rueckstiehs, 2025; Truong et al., 2025; Dao et al., 2026; Akram et al., 2026). Relational approaches match pairwise similarities, distances and angles, conditional affinities, graph edges, or topological summaries (Tung & Mori, 2019; Park et al., 2019; Passalis et al., 2021; Liu et al., 2019; Kim et al., 2024). Full Kernel Matrix Transfer uses Nystrom landmarks (Qian et al., 2022), and BINGO groups teacher-similar instances (Xu et al., 2022). Nevertheless, many relational objectives instantiate their relations from the current mini-batch; NRKD likewise mines an in-batch similarity matrix (Gou et al., 2025). This couples computational batching with geometric support selection.

2.2 BEYOND MINI-BATCH RELATIONAL SUPERVISION

Neighborhood graphs recover manifold structure in Laplacian Eigenmaps and Diffusion Maps, where a row-stochastic transition matrix over a k NN graph serves as the primitive description of local geometry (Belkin & Niyogi, 2003; Coifman & Lafon, 2006); heat-based constructions summarize the same operator at multiple scales (Sun et al., 2009; Huguet et al., 2023). Geometric KD aligns heat kernels between teacher and student graphs (Yang et al., 2022a), and graph-to-MLP methods distill observed graph structure into graphless models (Zhang et al., 2022; Tian et al., 2025); both assume a graph given by the task. GGPKD instead constructs the graph from cached text embeddings, so the teacher’s own retrieval behavior defines which relations a standard text encoder is supervised on. We use the transition operator at a single hop rather than its diffused powers: the one-hop rows already provide a teacher-weighted support and a per-row bandwidth, and multi-hop targets are treated as an alternative the method does not require. Neither the graph nor the teacher is needed at inference.

3 GGPKD: GLOBAL GEOMETRY THROUGH LOCAL ALIGNMENT

GGPKD distills an embedding teacher into a smaller student by matching the teacher’s one-hop k NN transition rows. The method has only two method-specific hyperparameters: the graph width `graph.k` and the auxiliary-row weight `row.weight`. All other temperatures, weights, candidate sets, and supervised rows are derived from the teacher graph. The graph and teacher embeddings are cached before training and are not needed at inference.

3.1 TEACHER GRAPH

Let $\mathcal{D} = \{x_i\}_{i=1}^N$ be the deduplicated training corpus. The teacher encodes every text once, and its embeddings are ℓ_2 -normalized so that all similarities are cosines. For each node i , we retrieve the $k = \text{graph.k}$ most similar non-self nodes. If $s_i^{(q)}$ is the q -th largest retrieved cosine, the row

bandwidth is

$$\tau_i = \frac{s_i^{(1)} - s_i^{(k)}}{\log k}. \quad (1)$$

Thus the k -th retrieved neighbor is k times less likely than the nearest neighbor. Because the bandwidth scales with any affine rescaling of the similarities, the resulting softmax row is affine-invariant and requires no teacher-specific temperature tuning.

The neighbor set is the retrieved list itself, $\mathcal{N}_i = \text{top-}k(i)$: we apply no reciprocity filter, so every relation the teacher retrieved for i is eligible to supervise i , and every node has degree exactly k . A mutual filter, which keeps $i \rightarrow j$ only when the two nodes retrieve each other, suppresses hubs but discards teacher-selected relations; we report it as an ablation in Section 4. The teacher transition row is

$$P_i^T(j) = \frac{\mathbf{1}[j \in \mathcal{N}_i] \exp(c_{ij}^T/\tau_i)}{\sum_{u \in \mathcal{N}_i} \exp(c_{iu}^T/\tau_i)}. \quad (2)$$

Because no edges are removed, the bandwidth in Equation (1) is read from the same scores that define the row, and from the same number of neighbors for every node.

3.2 DETERMINISTIC CANDIDATE POOL

The candidate set for anchor i is exactly its neighbor set $\mathcal{N}_i$. Every column the teacher retrieved is supervised and none is discarded, so no truncation constant enters the objective and every anchor is scored on the same number of columns, k . GGPKD draws no hard or random negatives and performs no per-epoch sampling: the candidates are a deterministic function of the graph and remain fixed throughout training. Within a mini-batch B , candidates are deduplicated into the shared pool

$$\mathcal{P}_B = \bigcup_{i \in B} \mathcal{N}_i, \quad (3)$$

so each distinct candidate text is encoded once.

3.3 TRAINING OBJECTIVE

The graph-scale loss matches each teacher row with a student softmax over the same columns and at the same stored bandwidth:

$$\mathcal{L}_{r=1} = \frac{1}{|B|} \sum_{i \in B} D_{\text{KL}}(P_i^T \parallel \text{softmax}_{j \in \mathcal{N}_i}(c_{ij}^S/\tau_i)). \quad (4)$$

This local row ranks candidates within an anchor’s own neighborhood. To also calibrate similarities across the batch, the ambient loss uses the complete shared pool. Its temperature is derived from the graph as $\tau_0 = \text{median}_i \tau_i$ and is shared by teacher and student:

$$\mathcal{L}_{r=0} = \frac{1}{|B|} \sum_{i \in B} D_{\text{KL}}(\text{softmax}_{j \in \mathcal{P}_B \setminus \{i\}}(c_{ij}^T/\tau_0) \parallel \text{softmax}_{j \in \mathcal{P}_B \setminus \{i\}}(c_{ij}^S/\tau_0)). \quad (5)$$

The ambient and graph groups receive equal total weight; this balance is fixed, not tuned.

GGPKD also reuses non-anchor candidates as auxiliary row centers. Let $\mathcal{J}_B = \mathcal{P}_B \setminus B$ and expose the portion of row j that is already present in the shared pool:

$$\Omega_j^B = \mathcal{N}_j \cap \mathcal{P}_B, \quad \mathcal{J}_B^+ = \{j \in \mathcal{J}_B : |\Omega_j^B| \geq 2\}. \quad (6)$$

Batch anchors are excluded because $\mathcal{L}_{r=1}$ already supervises their transition rows. The remaining rows are matched uniformly at their stored bandwidths:

$$\mathcal{L}_{\text{row}} = \frac{1}{|\mathcal{J}_B^+|} \sum_{j \in \mathcal{J}_B^+} D_{\text{KL}}(P_j^T|_{\Omega_j^B} \parallel \text{softmax}_{u \in \Omega_j^B}(c_{ju}^S/\tau_j)), \quad (7)$$

where the restricted teacher row is renormalized on Ω_j^B . This term requires no additional student encoding because all centers and columns are already in $\mathcal{P}_B$.

The complete objective is

$$\mathcal{L}_{\text{GGPKD}} = \underbrace{\mathcal{L}_{r=0} + \mathcal{L}_{r=1}}_{\mathcal{L}_{\text{rel}}} + \lambda_{\text{row}} \mathcal{L}_{\text{row}}. \quad (8)$$

The defaults are `graph_k` = 200 and $\lambda_{\text{row}} = \text{row_weight} = 1$. Increasing `graph_k` changes both neighborhood width and the bandwidth in Equation (1); those effects cannot be separated by a single-axis sweep.

Cost. Exact graph construction requires $O(N^2d)$ similarity work, although an approximate nearest-neighbor backend can replace the all-pairs search. During training, the dominant cost is encoding the deduplicated shared pool; $\mathcal{L}_{\text{row}}$ reuses those embeddings. At deployment, inference uses only the student encoder.

4 EXPERIMENTS

Table 1: Results across three teacher $\rightarrow$ student settings on nine datasets. Emotion, WiC, and STS-B (marked with $\star$) are in-domain; the remaining datasets are out-of-domain. Bold and underlined values denote the best and second-best results within each setting, excluding the Teacher and Student base reference rows. TALAS and GGPKD results are reported as mean $\pm$ standard deviation; the remaining baseline results are from TALAS (Dao et al., 2026). For setting (c), the epoch checkpoint is selected on the validation average.

	Classification (F1)			Pair Classification (AP)			STS (Spearman)			Domain Avg		Avg.
	Banking77	Tweet	Emotion $\star$	MRPC	SciTail	WiC $\star$	SICK	STS12	STS-B $\star$	Avg-In	Avg-Out	Avg.
(a) Qwen3-Embedding-4B $\rightarrow$ BERT-base (109M)												
Teacher	93.63	72.22	68.94	84.79	91.77	66.80	83.43	82.55	86.87	74.20	84.73	81.22
Student base	84.86	47.79	68.02	77.36	67.74	59.22	42.43	21.54	20.30	42.44	60.33	54.36
SimCSE-unsup	89.21	69.13	53.55	82.86	76.59	71.64	68.19	61.68	70.36	65.18	74.61	71.47
CDM	89.75	70.58	53.11	84.83	75.19	71.06	69.37	68.48	73.92	66.03	76.37	72.92
DSKD	89.66	69.93	54.51	83.13	77.95	71.53	68.65	63.37	71.42	65.82	75.45	72.24
Jasper and Stella	89.71	<u>73.83</u>	65.98	<u>83.33</u>	80.99	70.50	73.96	67.90	75.65	70.71	78.62	75.98
DistillCSE	<u>90.94</u>	70.90	60.80	82.86	78.98	68.63	75.12	72.61	77.08	68.84	78.57	75.32
EMO	90.78	73.54	67.48	84.46	81.57	69.76	75.66	68.83	77.17	71.47	79.14	76.58
TALAS	90.44 $\pm$ 0.13	74.62 $\pm$ 0.17	70.87 $\pm$ 0.52	85.90 $\pm$ 0.10	81.78 $\pm$ 0.29	68.84 $\pm$ 0.15	77.56 $\pm$ 0.10	73.12 $\pm$ 0.30	79.24 $\pm$ 0.10	72.98 $\pm$ 0.19	80.57 $\pm$ 0.01	78.04 $\pm$ 0.07
GGPKD	91.98 $\pm$ 0.09	73.74 $\pm$ 0.26	<u>68.78 $\pm$ 0.30</u>	84.22 $\pm$ 0.13	85.30 $\pm$ 0.16	66.28 $\pm$ 0.03	77.50 $\pm$ 0.06	75.95 $\pm$ 0.16	82.15 $\pm$ 0.07	72.40 $\pm$ 0.10	81.45 $\pm$ 0.12	78.43 $\pm$ 0.10
(b) BGE-M3 $\rightarrow$ MiniLMv2-H768 (66M)												
Teacher	93.52	73.85	68.56	85.81	91.87	61.47	79.18	78.73	84.87	71.63	83.83	79.76
Student base	81.84	47.70	71.67	79.67	65.06	60.89	49.46	26.76	24.12	44.24	62.41	56.35
SimCSE-unsup	87.89	71.45	55.65	83.87	76.32	67.96	66.93	59.66	63.97	62.53	74.35	70.41
CDM	89.15	70.48	58.71	84.16	75.91	70.00	68.14	62.74	67.78	65.50	75.10	71.90
DSKD	89.51	70.53	57.27	84.39	75.36	<u>69.96</u>	68.81	66.78	70.05	65.76	75.90	72.52
Jasper and Stella	88.38	74.13	62.81	<u>85.98</u>	80.33	68.38	74.35	66.80	74.55	68.58	78.33	75.08
DistillCSE	90.50	72.70	59.88	84.39	78.48	69.49	73.97	70.00	73.73	67.70	78.34	74.79
EMO	90.19	<u>74.26</u>	61.48	85.31	80.82	68.30	75.41	67.06	76.35	68.71	78.84	75.46
TALAS	89.72 $\pm$ 0.19	75.38 $\pm$ 0.11	65.83 $\pm$ 0.58	87.58 $\pm$ 0.02	82.90 $\pm$ 0.12	66.96 $\pm$ 0.07	75.82 $\pm$ 0.11	68.11 $\pm$ 0.20	77.51 $\pm$ 0.18	70.10 $\pm$ 0.13	79.92 $\pm$ 0.05	76.64 $\pm$ 0.04
GGPKD	90.78 $\pm$ 0.09	73.57 $\pm$ 0.32	66.45 $\pm$ 0.20	85.13 $\pm$ 0.06	83.50 $\pm$ 0.20	62.75 $\pm$ 0.20	78.27 $\pm$ 0.06	74.94 $\pm$ 0.18	81.74 $\pm$ 0.07	70.31 $\pm$ 0.03	81.03 $\pm$ 0.08	77.46 $\pm$ 0.05
(c) Qwen3-Embedding-0.6B $\rightarrow$ MiniLMv2-H384 (22M)												
Teacher	93.21	71.33	66.53	83.37	90.00	66.85	80.64	77.11	84.60	72.66	82.61	79.29
Student base	68.29	41.07	69.90	78.39	64.56	59.58	47.60	21.95	22.08	40.91	58.45	52.60
SimCSE-unsup	86.67	69.56	54.00	82.29	73.76	65.71	66.04	59.11	62.08	60.60	72.91	68.80
CDM	86.07	70.32	53.04	81.97	72.97	68.34	65.97	61.12	65.75	62.38	73.07	69.51
DSKD	85.66	69.86	52.25	82.17	73.52	<u>68.58</u>	66.54	62.52	66.70	62.51	73.38	69.76
Jasper and Stella	85.94	71.25	57.99	83.36	76.17	68.25	70.57	63.81	70.11	65.45	75.18	71.94
DistillCSE	88.03	69.87	54.24	83.20	77.66	67.88	70.09	68.19	71.53	64.55	76.17	72.30
EMO	87.03	<u>71.67</u>	59.63	83.45	78.17	68.89	<u>70.81</u>	65.78	71.42	<u>66.65</u>	76.15	72.98
TALAS	84.16 $\pm$ 0.15	72.82 $\pm$ 0.11	61.13 $\pm$ 0.52	84.85 $\pm$ 0.10	80.91 $\pm$ 0.27	66.32 $\pm$ 0.07	70.72 $\pm$ 0.15	67.80 $\pm$ 0.08	72.39 $\pm$ 0.31	66.61 $\pm$ 0.24	76.88 $\pm$ 0.07	73.46 $\pm$ 0.10
GGPKD	88.18 $\pm$ 0.03	70.78 $\pm$ 0.17	62.69 $\pm$ 0.14	<u>83.71 $\pm$ 0.07</u>	82.77 $\pm$ 0.11	65.39 $\pm$ 0.01	73.90 $\pm$ 0.08	74.57 $\pm$ 0.17	78.22 $\pm$ 0.02	68.77 $\pm$ 0.04	78.99 $\pm$ 0.05	75.58 $\pm$ 0.03

4.1 SENSITIVITY ANALYSIS

We evaluate GGPKD sensitivity on the Qwen3-Embedding-0.6B→MiniLMv2-H384 setting using seeds 42, 43, and 44. Tables 2 and 3 report downstream quality together with the measured resource cost. Increasing `graph_k` from 50 to 400 changes Avg. by only 0.34 points, but the largest setting costs substantially more time and GPU memory. Relative to the default `graph_k` = 200, the 400-neighbor setting gains 0.05 points while increasing mean training time by 36.8% and peak GPU memory by 54.2%. The row-loss sweep is also flat: Avg. spans 0.18 points, and the small timing differences are not sufficient to establish a causal cost effect. We therefore retain `graph_k` = 200 and $\lambda_{\text{row}} = 1$ as defaults rather than selecting a new setting on the test results.

5 CONCLUSION

REFERENCES

- Mohammad Kalim Akram, Saba Sturua, Nastia Havriushenko, Quentin Herreros, Michael Günther, Maximilian Werk, and Han Xiao. jina-embeddings-v5-text: Task-targeted embedding distillation. *CoRR*, abs/2602.15547, 2026. doi: 10.48550/arXiv.2602.15547. URL <https://arxiv.org/abs/2602.15547>.
- Anshumann, Mohd Abbas Zaidi, Akhil Kedia, Jinwoo Ahn, Taehwak Kwon, Kangwook Lee, Haejun Lee, and Joohyung Lee. Sparse logit sampling: Accelerating knowledge distillation in LLMs. In Wanxiang Che, Joyce Nabende, Ekaterina Shutova, and Mohammad Taher Pilehvar (eds.), *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 18085–18108, Vienna, Austria, July 2025. Association for Computational Linguistics. doi: 10.18653/v1/2025.acl-long.885. URL <https://aclanthology.org/2025.acl-long.885/>.
- Mikhail Belkin and Partha Niyogi. Laplacian eigenmaps for dimensionality reduction and data representation. *Neural Computation*, 15(6):1373–1396, 2003. doi: 10.1162/089976603321780317.
- Ronald R. Coifman and Stéphane Lafon. Diffusion maps. *Applied and Computational Harmonic Analysis*, 21(1):5–30, 2006. doi: 10.1016/j.acha.2006.04.006.
- Quoc Phong Dao, Hoang Son Nguyen, Khanh Chi Pham, Linh Ngo Van, Diep Thi-Ngoc Nguyen, Thien Huu Nguyen, and Trung Le. TALAS: Teacher-anchored layer alignment with adaptive sharpness-aware minimization for embedding distillation, 2026. Manuscript.
- Chaochen Gao, Xing Wu, Peng Wang, Jue Wang, Liangjun Zang, Zhongyuan Wang, and Songlin Hu. DistilCSE: Effective knowledge distillation for contrastive sentence embeddings. *CoRR*, abs/2112.05638, 2021. doi: 10.48550/arXiv.2112.05638. URL <https://arxiv.org/abs/2112.05638>.
- Yunfan Gao, Yun Xiong, Xinyu Gao, Kangxiang Jia, Jinliu Pan, Yuxi Bi, Yi Dai, Jiawei Sun, Meng Wang, and Haofen Wang. Retrieval-augmented generation for large language models: A survey. *arXiv preprint arXiv:2312.10997*, 2023.
- Jianping Gou, Xiaomeng Xin, Baosheng Yu, Heping Song, Weiyong Zhang, and Shaohua Wan. Neighborhood relation-based knowledge distillation for image classification. *Neural Networks*, 188:107429, 2025.
- Geoffrey E. Hinton, Oriol Vinyals, and Jeffrey Dean. Distilling the knowledge in a neural network. *CoRR*, abs/1503.02531, 2015. URL <http://arxiv.org/abs/1503.02531>.
- Guillaume Hugué, Alexander Tong, Edward De Brouwer, Yanlei Zhang, Guy Wolf, Ian Adelstein, and Smita Krishnaswamy. A heat diffusion perspective on geodesic preserving dimensionality reduction. In *Advances in Neural Information Processing Systems*, volume 36, 2023.
- Xiaoqi Jiao, Yichun Yin, Lifeng Shang, Xin Jiang, Xiao Chen, Linlin Li, Fang Wang, and Qun Liu. TinyBERT: Distilling BERT for natural language understanding. In Trevor Cohn, Yulan He, and Yang Liu (eds.), *Findings of the Association for Computational Linguistics: EMNLP*

Table 2: Sensitivity to the teacher-graph neighborhood size for Qwen3-Embedding-0.6B→MiniLMv2-H384. Downstream scores and training wall time are mean $\pm$ sample standard deviation over seeds 42, 43, and 44. Training-memory columns report the maximum per-run peak across those seeds; graph-build measurements contain one build per setting. Host and GPU memory are reported in GiB.

graph_k	Downstream quality				Training cost			Graph-build cost		
	Avg-In	Avg-Out	Avg. $\uparrow$	Δ Avg.	Wall (s)	Host	GPU	Wall (s)	Host	GPU
50	68.29 $\pm$ 0.09	78.82 $\pm$ 0.03	75.31 $\pm$ 0.04	-0.29	231.53 $\pm$ 16.35	9.82	7.98	36.8	2.19	6.44
100	68.46 $\pm$ 0.22	79.02 $\pm$ 0.09	75.50 $\pm$ 0.07	-0.10	257.23 $\pm$ 35.31	9.80	9.89	20.9	1.56	0.96
200 (default)	68.79 $\pm$ 0.07	79.00 $\pm$ 0.03	75.60 $\pm$ 0.04	ref.	329.90 $\pm$ 8.68	10.15	11.50	21.8	1.63	0.96
400	68.89 $\pm$ 0.14	79.02 $\pm$ 0.00	75.64 $\pm$ 0.04	+0.05	451.30 $\pm$ 18.93	10.59	17.74	28.6	2.11	0.98

Table 3: Sensitivity to the auxiliary row-relation weight for Qwen3-Embedding-0.6B→MiniLMv2-H384 at the default graph_k = 200. Downstream scores and wall time are mean $\pm$ sample standard deviation over seeds 42, 43, and 44. Host and GPU columns report the maximum observed per-run peak in GiB.

λ_{row}	Downstream quality				Training cost		
	Avg-In	Avg-Out	Avg. $\uparrow$	Δ Avg.	Wall (s)	Host	GPU
0	68.63 $\pm$ 0.15	78.86 $\pm$ 0.03	75.45 $\pm$ 0.04	-0.15	326.00 $\pm$ 10.06	10.08	11.97
0.25	68.73 $\pm$ 0.19	78.91 $\pm$ 0.06	75.52 $\pm$ 0.07	-0.08	331.53 $\pm$ 10.78	10.15	11.50
0.5	68.64 $\pm$ 0.11	78.99 $\pm$ 0.10	75.54 $\pm$ 0.03	-0.06	350.43 $\pm$ 19.74	10.14	11.50
1.0 (default)	68.79 $\pm$ 0.07	79.00 $\pm$ 0.03	75.60 $\pm$ 0.04	ref.	329.90 $\pm$ 8.68	10.15	11.50
2.0	68.91 $\pm$ 0.14	78.99 $\pm$ 0.01	75.63 $\pm$ 0.05	+0.03	301.23 $\pm$ 33.81	10.17	11.50

2020, pp. 4163–4174, Online, November 2020. Association for Computational Linguistics. doi: 10.18653/v1/2020.findings-emnlp.372. URL <https://aclanthology.org/2020.findings-emnlp.372/>.

Hee-Jun Jung, Doyeon Kim, Seung-Hoon Na, and Kangil Kim. Feature structure distillation with centered kernel alignment in bert transferring. *Expert Systems with Applications*, 234:120980, 2023. ISSN 0957-4174. doi: <https://doi.org/10.1016/j.eswa.2023.120980>. URL <https://www.sciencedirect.com/science/article/pii/S0957417423014823>.

Jungeun Kim, Junwon You, Dongjin Lee, Ha Young Kim, and Jae-Hun Jung. Do topological characteristics help in knowledge distillation? In *Proceedings of the 41st International Conference on Machine Learning*, volume 235, pp. 24674–24693. PMLR, 2024. URL <https://proceedings.mlr.press/v235/kim24aj.html>.

Seungyeon Kim, Ankit Singh Rawat, Manzil Zaheer, Sadeep Jayasumana, Veeranjaneyulu Sadhanala, Wittawat Jitkittum, Aditya Krishna Menon, Rob Fergus, and Sanjiv Kumar. Embeddistill: A geometric knowledge distillation for information retrieval. *arXiv preprint arXiv:2301.12005*, 2023.

Jinhyuk Lee, Zhuyun Dai, Xiaoqi Ren, Blair Chen, Daniel Cer, Jeremy R. Cole, Kai Hui, Michael Boratko, Rajvi Kapadia, Wen Ding, Yi Luan, Sai Meher Karthik Duddu, Gustavo Hernandez Abrego, Weiqiang Shi, Nithi Gupta, Aditya Kusupati, Prateek Jain, Siddhartha Reddy Jonnalagadda, Ming-Wei Chang, and Iftexhar Naim. Gecko: Versatile text embeddings distilled from large language models. *CoRR*, abs/2403.20327, 2024. doi: 10.48550/arXiv.2403.20327. URL <https://arxiv.org/abs/2403.20327>.

Yufan Liu, Jiajiong Cao, Bing Li, Chunfeng Yuan, Weiming Hu, Yangxi Li, and Yunqiang Duan. Knowledge distillation via instance relationship graph. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 7096–7104, 2019.

Niklas Muennighoff, Nouamane Tazi, Loïc Magne, and Nils Reimers. MTEB: Massive text embedding benchmark. In *Proceedings of the 17th Conference of the European Chapter of the Association for Computational Linguistics*, pp. 2014–2037. Association for Computational Linguistics,

2023. doi: 10.18653/v1/2023.eacl-main.148. URL <https://aclanthology.org/2023.eacl-main.148/>.
- Wonpyo Park, Dongju Kim, Yan Lu, and Minsu Cho. Relational knowledge distillation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 3967–3976, 2019. URL <https://arxiv.org/abs/1904.05068>.
- Nikolaos Passalis, Maria Tzelepi, and Anastasios Tefas. Probabilistic knowledge transfer for lightweight deep representation learning. *IEEE Transactions on Neural Networks and Learning Systems*, 32(5):2030–2039, 2021. doi: 10.1109/TNNLS.2020.2995884.
- Qi Qian, Hao Li, and Juhua Hu. Improved knowledge distillation via full kernel matrix transfer. In *Proceedings of the 2022 SIAM International Conference on Data Mining (SDM)*, pp. 612–620, 2022. doi: 10.1137/1.9781611977172.69.
- Nils Reimers and Iryna Gurevych. Sentence-BERT: Sentence embeddings using Siamese BERT-networks. In Kentaro Inui, Jing Jiang, Vincent Ng, and Xiaojun Wan (eds.), *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pp. 3982–3992, Hong Kong, China, November 2019. Association for Computational Linguistics. doi: 10.18653/v1/D19-1410. URL <https://aclanthology.org/D19-1410/>.
- Victor Sanh, Lysandre Debut, Julien Chaumond, and Thomas Wolf. Distilbert, a distilled version of BERT: smaller, faster, cheaper and lighter. *CoRR*, abs/1910.01108, 2019. URL <http://arxiv.org/abs/1910.01108>.
- Jian Sun, Maks Ovsjanikov, and Leonidas J. Guibas. A concise and provably informative multi-scale signature based on heat diffusion. *Computer Graphics Forum (Proc. SGP)*, 28(5):1383–1392, 2009.
- Yijun Tian, Shichao Pei, Xiangliang Zhang, Chuxu Zhang, and Nitesh V. Chawla. Knowledge distillation on graphs: A survey. *ACM Comput. Surv.*, 57(8):189:1–189:16, 2025. doi: 10.1145/3711121. URL <https://doi.org/10.1145/3711121>.
- Minh-Phuc Truong, Hai An Vu, Tu Vu, Nguyen Thi Ngoc Diep, Linh Ngo Van, Thien Huu Nguyen, and Trung Le. EMO: Embedding model distillation via intra-model relation and optimal transport alignments. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pp. 7594–7606. Association for Computational Linguistics, 2025. doi: 10.18653/v1/2025.emnlp-main.385. URL <https://aclanthology.org/2025.emnlp-main.385/>.
- Frederick Tung and Greg Mori. Similarity-preserving knowledge distillation. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pp. 1365–1374, 2019.
- Robin Vujanic and Thomas Rückstieß. LEAF: Knowledge distillation of text embedding models with teacher-aligned representations. In Maria Liakata, Viviane P. Moreira, Jiajun Zhang, and David Jurgens (eds.), *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 43362–43383, San Diego, California, United States, July 2026. Association for Computational Linguistics. ISBN 979-8-89176-390-6. doi: 10.18653/v1/2026.acl-long.2008. URL <https://aclanthology.org/2026.acl-long.2008/>.
- Robin Vujanic and Thomas Rueckstieß. LEAF: Knowledge distillation of text embedding models with teacher-aligned representations. *CoRR*, abs/2509.12539, 2025. doi: 10.48550/arXiv.2509.12539. URL <https://arxiv.org/abs/2509.12539>.
- Wenhui Wang, Furu Wei, Li Dong, Hangbo Bao, Nan Yang, and Ming Zhou. Minilm: Deep self-attention distillation for task-agnostic compression of pre-trained transformers. In Hugo Larochelle, Marc’Aurelio Ranzato, Raia Hadsell, Maria-Florina Balcan, and Hsuan-Tien Lin (eds.), *Advances in Neural Information Processing Systems 33: Annual Conference on Neural Information Processing Systems 2020, NeurIPS 2020, December 6-12, 2020, virtual*, 2020. URL <https://proceedings.neurips.cc/paper/2020/hash/3f5ee243547dee91fbd053c1c4a845aa-Abstract.html>.

- Haohang Xu, Jiemin Fang, Xiaopeng Zhang, Lingxi Xie, Xinggang Wang, Wenrui Dai, Hongkai Xiong, and Qi Tian. Bag of instances aggregation boosts self-supervised distillation. In *International Conference on Learning Representations*, 2022.
- Jiahao Xu, Wei Shao, Lihui Chen, and Lemao Liu. DistillCSE: Distilled contrastive learning for sentence embeddings. *CoRR*, abs/2310.13499, 2023. doi: 10.48550/arXiv.2310.13499. URL <https://arxiv.org/abs/2310.13499>.
- Xiaohan Xu, Ming Li, Chongyang Tao, Tao Shen, Reynold Cheng, Jinyang Li, Can Xu, Dacheng Tao, and Tianyi Zhou. A survey on knowledge distillation of large language models. *CoRR*, abs/2402.13116, 2024. doi: 10.48550/arXiv.2402.13116. URL <https://arxiv.org/abs/2402.13116>.
- Chenxiao Yang, Qitian Wu, and Junchi Yan. Geometric knowledge distillation: Topology compression for graph neural networks. In *Advances in Neural Information Processing Systems*, volume 35, pp. 29761–29775, 2022a.
- Chuanguang Yang, Helong Zhou, Zhulin An, Xue Jiang, Yongjun Xu, and Qian Zhang. Cross-image relational knowledge distillation for semantic segmentation. In *2022 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 12309–12318, 2022b. doi: 10.1109/CVPR52688.2022.01200.
- Dun Zhang, Jiacheng Li, Ziyang Zeng, and Fulong Wang. Jasper and stella: Distillation of SOTA embedding models. *CoRR*, abs/2412.19048, 2024. doi: 10.48550/arXiv.2412.19048. URL <https://arxiv.org/abs/2412.19048>.
- Shichang Zhang, Yozen Liu, Yizhou Sun, and Neil Shah. Graph-less neural networks: Teaching old mlps new tricks via distillation. In *The Tenth International Conference on Learning Representations, ICLR 2022, Virtual Event, April 25-29, 2022*. OpenReview.net, 2022. URL https://openreview.net/forum?id=4p6_5HBWPCw.
- Yanzhao Zhang, Mingxin Li, Dingkun Long, Xin Zhang, Huan Lin, Baosong Yang, Pengjun Xie, An Yang, Dayiheng Liu, Junyang Lin, Fei Huang, and Jingren Zhou. Qwen3 embedding: Advancing text embedding and reranking through foundation models. *CoRR*, abs/2506.05176, 2025. doi: 10.48550/arXiv.2506.05176. URL <https://arxiv.org/abs/2506.05176>.